



FIT1043 Lecture 8

Introduction to Data Science

Mahsa Salehi


Faculty of Information Technology, Monash University

Semester 2, 2022

Unit Schedule

Week	Activities	Assignments
1	Overview of data science	Weekly Lecture/tutorial active participation assessment
2	Introduction to Python for data science	
3	Data visualisation and descriptive statistics	
4	Data sources and data wrangling	
5	Data analysis theory	Assignment 1
6	Regression analysis	
7	Classification and clustering	
8	Introduction to R for data science	
9	Characterising data and "big" data	Assignment 2
10	Big data processing	
11	Issues in data management	
12	Industry guest lecture	Assignment 3

Our Standard Value Chain



This Week

Tools for
data science

Outline

- Motivation to study R
- R data types
- Essential libraries
 - Wrangling
 - Exploration and analysis
 - Visualisation

Learning Outcomes (Week 8)


By the end of this week you should be able to:

- Comprehend essentials for coding in R for data science
- Explain and interpret given R commands
- Apply R commands for data wrangling, visualisation, exploration and analysis

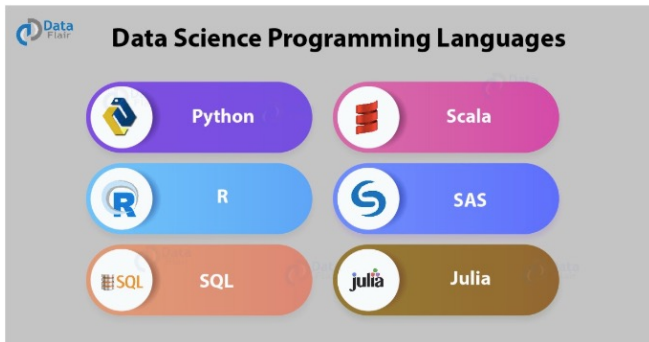


Introduction to R for Data Science

[R for Data Science](#) by H. Wickham
and G. Grolemund



Top 6 Data Science Programming Languages for 2019



What is R?



- Data science preferred tools
- A language for **analysing** and **visualising** data
 - Interpreted (scripting) language, so no need to compile code
 - Designed by statisticians
 - Open-source
 - Very popular!




image src: kdnuggets.com

FLUX Question

- R Or Python



R Vs Python



Parameter	R	Python
Objective	Data analysis and statistics	Deployment and production
Flexibility	Easy to use available library	Easy to construct new models from scratch.
Important Packages and library	tidyverse, ggplot2, caret, zoo	pandas, scipy, scikit-learn, TensorFlow, caret
Disadvantages	Steep Learning curve Dependencies between library	Not as many specialized packages for statistical computing as R
Comparison	<ul style="list-style-type: none">• Functional• More statistical support in general	<ul style="list-style-type: none">• Object Oriented• More straightforward to do non-statistical tasks

Setting Up R Environment

Installing R:

- Available for download from the R project
 - <https://www.r-project.org/>
- Get the **Rstudio IDE** (Integrated Development Environment) from:
 - <https://www.rstudio.com/products/rstudio/>
 - Both open source and commercial versions

Running R:

- Either type “R” in a shell (Linux/MacOS)
- Or start R-Studio application (Windows/MacOS)

R Environment

```
dyn-130-194-70-208:scripts mcarmans$ R

R version 3.3.0 (2016-05-03) -- "Supposedly Educational"
Copyright (C) 2016 The R Foundation for Statistical Computing
Platform: x86_64-apple-darwin13.4.0 (64-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

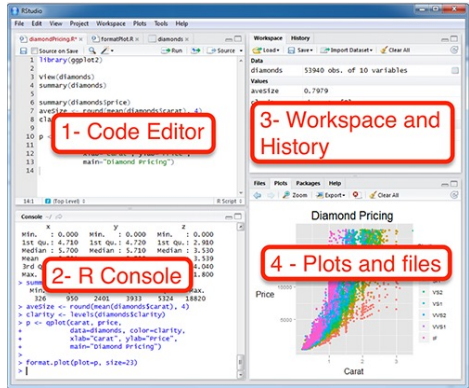
Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

[Previously saved workspace restored]

> 
```



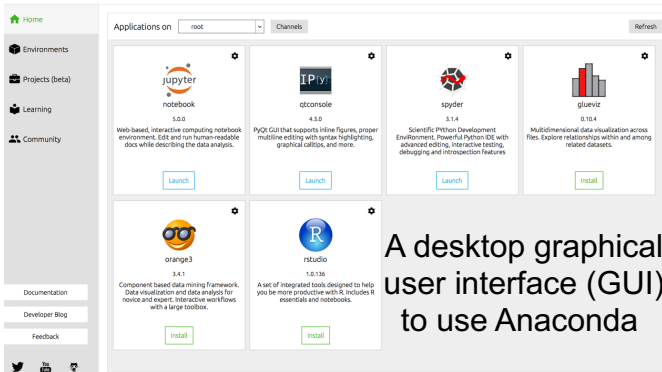
Anaconda Project

- [Anaconda](#) is a package manager, an environment manager, a Python/R data science distribution, and a collection of over 1,500+ open source packages. Anaconda is free and easy to install.



ANACONDA NAVIGATOR

Sign in to Anaconda Cloud

The screenshot shows the Anaconda Navigator desktop application. On the left is a sidebar with navigation links: Home, Environments, Projects (beta), Learning, and Community. Below these are links for Documentation, Developer Blog, and Feedback, along with social media icons for Twitter, YouTube, and GitHub. The main area displays a grid of application tiles. The top row includes Jupyter Notebook (5.0.0), QtConsole (4.3.0), Spyder (3.1.4), and Glueviz (0.10.4). The bottom row includes Orange3 (3.4.1) and RStudio (1.0.136). Each tile shows the application's icon, name, version, a brief description, and a button to either 'Launch' or 'Install' the application. At the top of the main area, there are dropdown menus for 'Applications on' (set to 'root') and 'Channels', and a 'Refresh' button.

A desktop graphical user interface (GUI) to use Anaconda

Built-in Data Sets in R

- List available data set :
`> data()`
- Load a built-in dataset
`> data(mtcars)`
- Inspect the data set
`> head(mtcars, 6)`


	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Basic R Syntax

- Compute mathematical expressions:

```
> 2^3+2
```

```
[1] 10
```




Here **>** denotes the command prompt & the output is prefixed by: **[1]**

- Define variables and assign values:

```
> A <- 10
```

```
> 5*A +6
```

```
[1] 56
```



It is traditional in R to use left-arrow for assignment, but you can also use equals:

```
> A = 10
```

Operators

- Arithmetic Operators

Operator	Description
+	addition
-	subtraction
*	multiplication
/	division
^ or **	exponentiation
x %% y	modulus (x mod y) 5%%2 is 1
x %/% y	integer division 5%/%2 is 2

- Logical Operators

Operator	Description
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	exactly equal to
!=	not equal to
!x	Not x
x y	x OR y
x & y	x AND y
isTRUE(x)	test if X is TRUE

If Else Condition

Syntax: statement would be executed if expression is TRUE

```
if(expression)
{
  statement/s
}
```

Example:

```
> x <- 10
> if(x>0)
{
  print("This is Positive Number)
}
[1] "This is Positive number"
```

For Loop

Syntax: statement would be executed n-times.

```
for(i in 1:n)
{
statement/s
}
```

Example:

```
> for(i in 1:3)
{
  print(i^2)
}
[1] 1
[1] 4
[1] 9
```

While Loop

Syntax:

```
while(condition)
{
  statement/s
}
```

Example:

```
> i <- 1
> while (i <=6) {
  print(i*i)
  i = i+1
}
[1] 1
[1] 4
[1] 9
```

Break Statement

- **Break:** Stop the iteration and exit the loop.

Example:

```
> x <- 1:5  
> for (i in x) {  
  if (i == 3){  
    break  
  }  
  print(i)  
}
```

```
[1] 1
```

```
[1] 2
```

Next Statement

- **Next:** Skip one step of the loop and jumps to the next cycle.

Example:

```
> x <- 1:5
```

```
> for (i in x) {  
  if (i == 3){  
    next  
  }  
  print(i)  
}
```

```
[1] 1
```

```
[1] 2
```

```
[1] 4
```

```
[1] 5
```

Basic Data Types

- Numeric

```
> x <- 10.5
```

- Integer

```
> x <- as.integer(10.5)
```

- Complex

```
> x <- 1+2i
```

- Logical

```
> x <- TRUE
```

- Character

```
> x <- "Intro To R"
```

Basic Data Types(Cont.)

- Print the class name of y

```
> y <- 8  
> class(y)  
[1] "numeric"
```

- Is y an integer?

```
> is.integer(y)  
[1] FALSE
```

- Change data type

```
> as.character(y)  
[1] "8"
```

- Getting help

```
> help(c)
```

What is Vector?

R has **c()** built in function which allows to store more than one value.

- Define a vector using the concatenate function:

```
> B <- c(5,6,3,0)
```

```
> B
```


```
[1] 5 6 3 0
```

- Concatenate function can be applied to vectors too:

```
> B <- c(B,c(1,2))
```

```
> B
```

```
[1] 5 6 3 0 1 2
```



You **must** use the concatenate function **c()** to build a vector, just writing **(5,6,3,0)** won't work!

Accessing Vector Elements

- Accessing vector elements using position


```
> x <- c("Jan","Feb","Mar","April")
```

```
> y <- x[c(1,3,4)]
```

```
> print(y)
```

```
[1] "Jan" "Mar" "April"
```

Unlike Python, the first element of an array has index 1 (not 0)



- Accessing vector elements using negative indexing

```
> t <- x[c(-1,-4)]
```

```
> print(t)
```


```
[1] "Feb" "Mar"
```

- Access range of values in vector

```
> x[1:3]
```

```
[1] "Jan" "Feb" "Mar"
```

The colon operator **1:n** generates a vector of integers from 1 to n, **inclusive**:



Vector Arithmetic Operation

Operations can be performed on two vectors (**same length**) directly and are interpreted in an element-wise fashion.

- Create two vectors.
 `> v1 <- c(1,2,4,5,7,11)`
 `> v2 <- c(12,4,3,8,1,21)`
- Vector multiplication.
 `> multi.result <- v1*v2`
 `> print(multi.result)`
 `[1] 12 8 12 40 7 231`

FLUX Question

- How to check if a vector contains missing values?



What is Data Frame?

We can combine vectors together to form a table, called a “data frame”

- Create the data frame

```
> names <- c("Bill", "Ted", "Henry", "Joan")  
> ages <- c(76, 82, 104, 78)  
> heights <- c(1.55, 1.69, 1.49, 1.57)  
> myTable <- data.frame(names, ages, heights)  
> print(myTable)
```

	names	ages	heights
--	-------	------	---------

1	Bill	76	1.55
2	Ted	82	1.69
3	Henry	104	1.49
4	Joan	78	1.57

Rename The Columns Of Data Frame

R has **names(df)** built in function which allows you to rename data frame columns

- Pass a vector of new names to the function
> names(myTable)<-c("Names", "Ages", "Heights")
> print(myTable)

	Names	Ages	Heights
1	Bill	76	1.55
2	Ted	82	1.69
3	Henry	104	1.49
4	Joan	78	1.57

Accessing Elements of Data Frame

- Number of rows in data frame
`> nrow(myTable)`
[1] 4
- Number of columns in data frame
`> ncol(myTable)`
[1] 3
- Dimension of data frame
`> dim(myTable)`
[1] 4 3

Get the Structure of the Data Frame

- Display the column names and data types
> `str(myTable)`

```
'data.frame': 4 obs. of 3 variables:  
 $ Names : chr "Bill" "Ted" "Henry" "Joan"  
 $ Ages : num 76 82 104 78  
 $ Heights: num 1.55 1.69 1.49 1.57
```

Summary Statistic

- Minimum value
`> min(myTable$Ages)`
[1] 76
- Average value
`> mean(myTable$Heights)`
[1] 1.575
- Standard deviation
`> sd(myTable$Heights)`
[1] 0.08386497

Summary of Data Frame

```
> summary(myTable)
```

Names	Ages	Heights
Length:4	Min. : 76.0	Min. :1.490
Class :character	1st Qu.: 77.5	1st Qu.:1.535
Mode :character	Median : 80.0	Median :1.560
	Mean : 85.0	Mean :1.575
	3rd Qu.: 87.5	3rd Qu.:1.600
	Max. :104.0	Max. :1.690

Extracting Data From Data Frame

- Accessing column/s by name
> `myTable["Ages"]`
- Accessing multiple columns by name
> `myTable[c("Names", "Ages")]`
- Accessing columns by index
> `myTable[2]`
- Accessing multiple columns by index
> `myTable[c(1,2)]`

Extracting Data From Data Frame(Cont.)

- Accessing first row and all the columns by appending comma

```
> myTable[1,]
```

	Names	Ages	Heights
1	Bill	76	1.55

Strange looking syntax for selecting rows is due to fact that in R, tables are matrices that are indexed by **[row,column]** (i.e. row first)

- Accessing a range of rows and all the columns

```
> myTable[2:4,]
```

	Names	Ages	Heights
2	Ted	82	1.69
3	Henry	104	1.49
4	Joan	78	1.57

Extracting Data From Data Frame(Cont.)

- Accessing particular cells by [row,column]
 > myTable[1,2]
 [1] 76
 > myTable[3:4,2:3]
 Ages Heights
 3 104 1.49
 4 78 1.57
- Referring to a variable (a column) by using the \$ syntax:
 > myTable\$Ages
 [1] 76 82 104 78
 > myTable\$Ages[3]
 [1] 104

Sorting Data in Data Frame

- Sort by Ages
`> newData <- myTable[order(myTable$Ages),]`
- Sort by Ages (ascending)
`> newData <- myTable[order(myTable$Ages,
decreasing = TRUE),]`
- Sort by Ages and Heights
`> newData <-
myTable[order(myTable$Ages,myTable$Heights),]`

Merging Data in Data Frame

merge(): Used to merge two data frames by common key variable/s

- Merge two data frames by ID
> total <- merge(myTable, myTable2, by="Names")

rbind() Used to join two data frames vertically(Must have same number of variables)

- Join two data frames
> total <- rbind(myTable, myTable2)

Getting & Setting Working Directory

Before reading/writing in R it is important to specify the location where we can find the respective file to read/write.

- Get the current working directory.
 `> getwd()`
 [1] " C:/Users/username/FolderName"
- Set current working directory.
 `> setwd(" D:/FolderName ")`

Writing CSV File

R has a **write.csv()** built in function to write data into a CSV file.

- Write a data into csv file (file is in current working directory)

```
> write.csv(myTable, "FileName.csv")
```

- Read a csv file (file is in other location)

```
> write.csv(myTable, "D:/FolderName/FileName.csv")
```

Reading CSV File

R has a **read.csv()** built in function to read a CSV file.

- Read a csv file (file is in current working directory)
 > myData = read.csv("FileName.csv")
 > print(myData)
- Read a csv file (file is in other location)
 > myData = read.csv("D:/FolderName/FileName.csv")
 > print(myData)

Displaying Data

- If a file is big, we don't want to print it all out, just to have a look at it. Instead we can inspect the first/last lines of the table:

```
> head(myData)
```

```
...
```

```
> tail(myData)
```


```
...
```

Loading Libraries

- Libraries are lists of functions that are not available in R by default.
- Loading a library

```
> library(moments)
```

```
> skewness(data)
```



The `skewness ()` function is provided by the `moments` library
- Before loading a library for the first time you will need to install the package on your machine:

```
> install.packages("moments")
```


Visualising Data: Bar Chart

- Compare the value for categorical data using bar chart

```
> H <- c(25,12,43,7,51)
```


```
> M <- c("Delhi","Beijing","Washington","Tokyo","Moscow")
```

```
> barplot(H,xlab="Month",ylab="Happiness Index", col="blue",  
          names.arg=M, main="Happiness Index",border="red")
```



Visualising Data: Group Bar Chart


```
> counts <- table(mtcars$vs, mtcars$gear)
> barplot(counts, main="Car Distribution by Gears and VS", xlab =
  "Number of Gears", col=c("darkblue","red"),legend =
  rownames(counts),beside=TRUE)
```



Visualising Data: Histogram

Inspect the distribution of values for a particular variable by plotting it as a histogram

```
> hist(mtcars$hp,main="Histogram for mtcars",xlab="HP",  
border="red",col="blue",xlim=c(0,400))
```



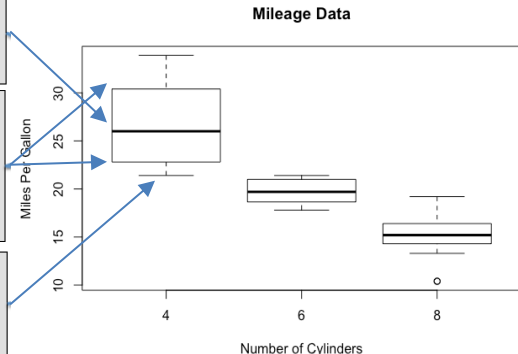
Visualising Data: Boxplot

- Or its summary statistics by plotting it as a boxplot
 - > `boxplot(mpg ~ cyl, data=mtcars, xlab="Number of Cylinders", ylab="Miles Per Gallon", main="Mileage Data")`

Median value (half the data lies above and the other half below)

Upper & lower quartiles (25% of the data lies above/below these values and 50% between them)

Minimum value (or $1.5 \times \text{InterQuartileRange}$ below lower quartile)



FLUX Question

How to find outliers in R?




Visualising Data: Scatter Plot

- Or the variation of one variable against another by plotting data as a scatterplot

```
> input <- mtcars[,c('wt','mpg')]
```

```
> plot(x=input$wt,y=input$mpg, xlab="Weight", ylab="Mileage",
      xlim=c(2.5,5), ylim=c(15,30), main="Weight vs Mileage")
```



Linear Regression

Often we'd like to see if there exists a linear trend relationship between two variables.

- Creating sample Data for height and weight

```
> height <- c(151, 174, 138, 186, 128, 136, 179, 163, 152, 131)
> weight <- c(63, 81, 56, 91, 47, 57, 76, 72, 62, 48)
```

- Fitting a linear model in R is very simple

```
> fit <- lm(height~weight)
> print(fit)
```

Call:

`lm(formula = height ~ weight)`

Coefficients:

(Intercept)	weight
61.380	1.415

Linear Regression(Cont.)

- Print out summary information regarding the fit (the slope, etc.)

```
> summary(fit)
```

Call:

```
lm(formula = height ~ weight)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.0529	-2.4833	-0.0912	1.3774	10.0562

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	61.3803	7.2653	8.448	2.94e-05 ***
weight	1.4153	0.1089	12.997	1.16e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.712 on 8 degrees of freedom


Multiple R-squared: 0.9548, Adjusted R-squared: 0.9491

F-statistic: 168.9 on 1 and 8 DF, p-value: 1.164e-06

Visualize the Regression

- Plot the chart

```
> plot(height, weight,col = "blue",main = "Height & Weight  
Regression", abline(lm(weight ~ height)),cex = 1.3,pch= 16,  
xlab = "Weight in Kg",ylab = "Height in cm")
```



Decision Tree

- Install and load the party package.
 > install.packages("party")
 > library(party)
- Create the input data frame
 > inputData <- readingSkills[c(1:105),]
 > print(inputData)

	nativeSpeaker	age	shoeSize	score
1	yes	5	24.83189	32.29385
2	yes	6	25.95238	36.63105
3	no	11	30.42170	49.60593
4	yes	7	28.66450	40.28456
5	yes	11	31.88207	55.46085
6	yes	10	30.07843	52.83124

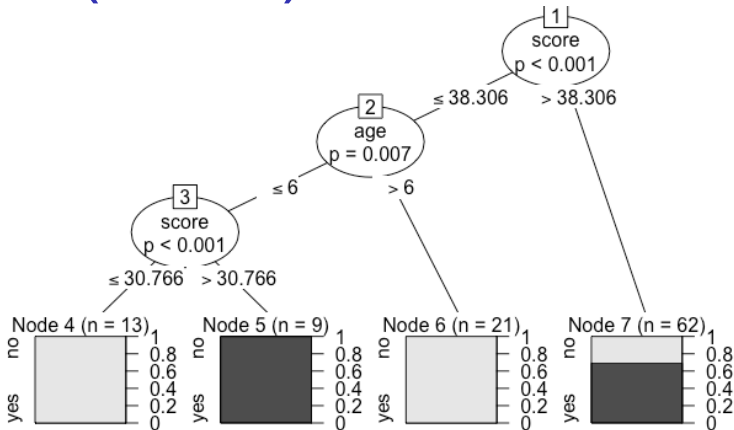
Visualize The Decision Tree

- Create the tree

```
> outputTree <- ctree( nativeSpeaker ~ age + shoeSize  
                        + score, data = inputData)
```
- Plot the tree

```
> plot(outputTree)
```

Visualize The Decision Tree(Cont.)



End Of Introduction

- You will work on a large data file in your tutorial/lab this week.
 - hourly ozone level readings across the US.
- There are MANY excellent R resources online if you'd like to learn more. For example:
 - [lynda.com](https://www.lynda.com)
 - [datacamp.com](https://www.datacamp.com)